



INTERNATIONAL
BRAIN
LABORATORY

Fiete Lab

2020 Trimester 1 Goals

Prof. Ila Fiete, Dr. Leenoy Meshulam, Rylan Schaeffer
January 28, 2020

Project 1 - Overview

See Leenoy's presentation from last week's Theory Meeting

[https://drive.google.com/drive/u/1/folders/
1IuUrh0hq1S18MBfCQkXGkLwz1OnKXDZ7](https://drive.google.com/drive/u/1/folders/1IuUrh0hq1S18MBfCQkXGkLwz1OnKXDZ7)

Project 2 - Overview

Goals:

1. Understand how neurally-plausible mechanistic models solve IBL task

Project 2 - Overview

Goals:

1. Understand how neurally-plausible mechanistic models solve IBL task
2. Leverage understanding to direct exploration/analysis of biological circuit data

Project 2 - Overview

Goals:

1. Understand how neurally-plausible mechanistic models solve IBL task
2. Leverage understanding to direct exploration/analysis of biological circuit data

Research Questions

Project 2 - Overview

Goals:

1. Understand how neurally-plausible mechanistic models solve IBL task
2. Leverage understanding to direct exploration/analysis of biological circuit data

Research Questions

- What computation(s) do plausible neural models execute to solve the same IBL task solved by mice?

Project 2 - Overview

Goals:

1. Understand how neurally-plausible mechanistic models solve IBL task
2. Leverage understanding to direct exploration/analysis of biological circuit data

Research Questions

- What computation(s) do plausible neural models execute to solve the same IBL task solved by mice?
- How do model design choices (e.g. learning rules, size, connectivity, noise) affect the learnt solution, and why?

Project 2 - Overview

Goals:

1. Understand how neurally-plausible mechanistic models solve IBL task
2. Leverage understanding to direct exploration/analysis of biological circuit data

Research Questions

- What computation(s) do plausible neural models execute to solve the same IBL task solved by mice?
- How do model design choices (e.g. learning rules, size, connectivity, noise) affect the learnt solution, and why?
- How does gradient descent on a high dimensional error landscape enable learning the IBL task's structure?

Project 2 - Implementation Details


- Loss: Cross Entropy
- Stimuli sampled from $\mathcal{N}(-1, 1)$ and $\mathcal{N}(1, 1)$ plus block prior
- Architectures: RNN (tanh), LSTM, GRU
- Number of stacked layers: 1
- Hidden dimension: 10, 20, 50, 100, 200
- Weight Initializations: Uniform $\pm \frac{1}{\text{hidden state size}}$, Identity
- Optimizer: SGD with LR=0.01, no momentum

Project 2 - Preliminary Results

Models appear to learn task:

Figure 1: Loss vs Grad Step

loss_per_grad_step
tag: train/loss_per_grad_step




Project 2 - Preliminary Results

Models appear to learn task:

Figure 2: $P(\text{Correct Choice})$ vs Grad Step


avg_correct_choice_per_grad_step
tag: train/avg_correct_choice_per_grad_step



Project 2 - Preliminary Results

Models appear to learn task's block structure:


Figure 3: $P(\text{Correct Action})$ vs Trial Number within Block



Project 2 - Preliminary Results

Models appear to learn task's block structure:


Figure 3: $P(\text{Correct Action})$ vs Trial Number within Block



Project 2 - Preliminary Results

Models appear to learn task's block structure:

Figure 3: $P(\text{Correct Action})$ vs Trial Number within Block



nn: num layers=1, hidden size=10, param init=eye


Project 2 - Preliminary Results

Architecture appears to effect learning strategy.

Consider $P(\text{Correct Action})$ after first indication that model recognizes block structure:


avg_model_prob_by_trial_num_within_block lstm, num_layers=1, hidden_size=10, param_init=default_2020-01-27 14:52:59.43377
step 280 Mon Jan 27 2020 18:05:05 Eastern Standard Time

Average Model Probability for Correct Choice by Trial within Block



avg_model_prob_by_trial_num_within_block lstm, num_layers=1, hidden_size=10, param_init=default_2020-01-27 14:52:59.44066
step 410 Mon Jan 27 2020 17:46:03 Eastern Standard Time

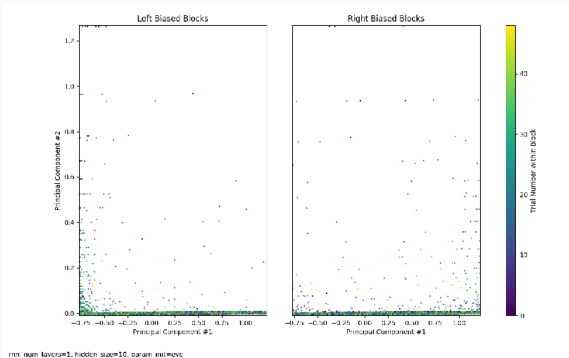
Average Model Probability for Correct Choice by Trial within Block



Project 2 - Preliminary Results

Models (sometime) reveal 2D manifold reflecting task structure:

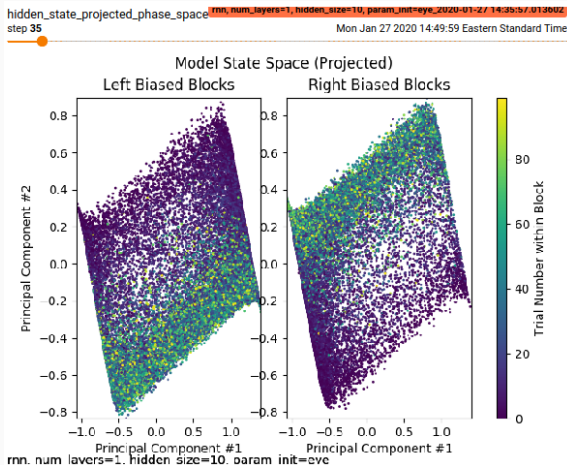
Figure 4: RNN Phase State projected onto first two PCs



Project 2 - Preliminary Results

Models (sometime) reveal 2D manifold reflecting task structure:


Figure 4: RNN Phase State projected onto first two PCs



Project 2 - Preliminary Results

Movement in projected phase space for left, neutral, right stimuli:

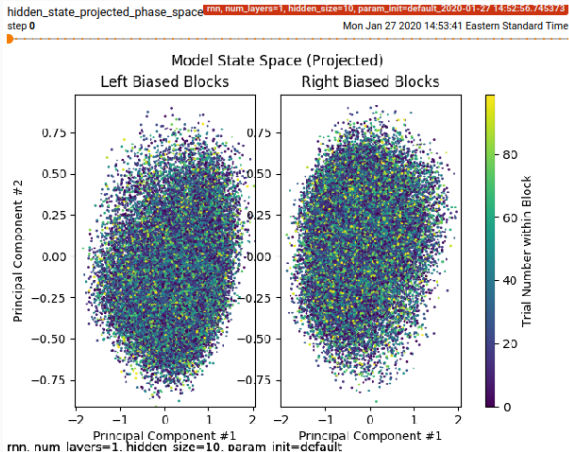
Figure 5: $P(\text{Correct Choice})$ vs Grad Step



Project 2 - Preliminary Results

Models (sometime) do not appear to learn 2D manifold:

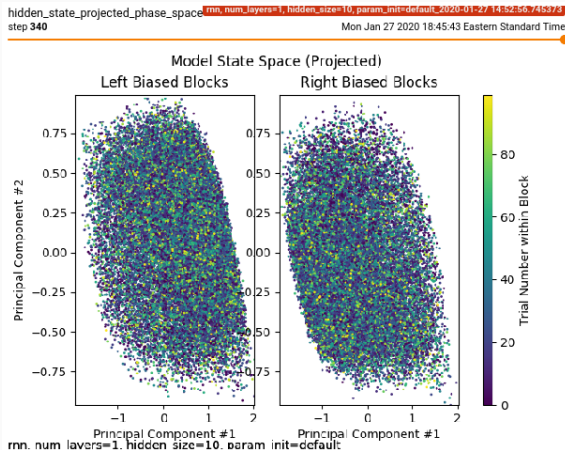
Figure 6: RNN Phase State projected onto first two PCs



Project 2 - Preliminary Results

Models (sometime) do not appear to learn 2D manifold:

Figure 6: RNN Phase State projected onto first two PCs



Project 2 - Preliminary Results

Figure 7: Hidden units correlations before learning.

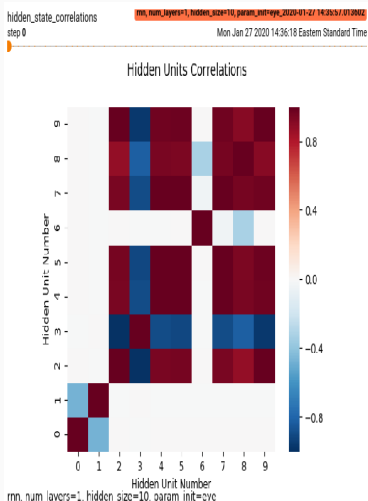
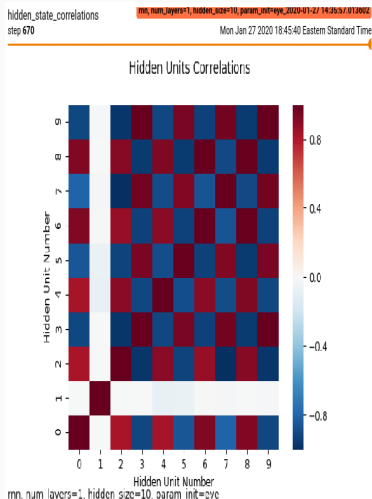



Figure 8: Hidden units correlations after learning.



Project 2 - Preliminary Results

Other manifolds,

Figure 9: RNN Phase State projected onto first two PCs



Project 2 - Preliminary Results

Other manifolds,

Figure 9: RNN Phase State projected onto first two PCs

